

Volunteer Cloud Computing

Dany Wilson – Dr. Stéphane Somé

University of Ottawa

January 23, 2015

Agenda

- 1 Introduction
- 2 Related Work
- 3 Infrastructure
- 4 Open Problems
- 5 Contributions
- 6 Future Work

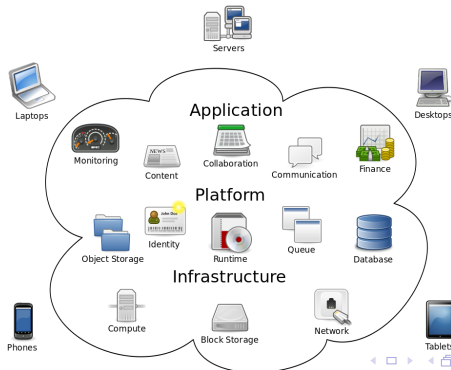
Volunteer Computing



- Great Internet Mersenne Prime Search [1996]
- Distributed Computing based on Collaboration
- ... throughput of 137.023 TeraFLOP/s

Cloud Computing

- Natural evolution of Web 2.0, SoA and Virtualization technologies.




Definition

NIST provided a description of the characteristics of a Cloud Computing infrastructure:

- On-demand Self-Service
- Broad Network Access
- Resource Pooling
- Rapid Elasticity
- Measured Services

Separation of Responsibilities



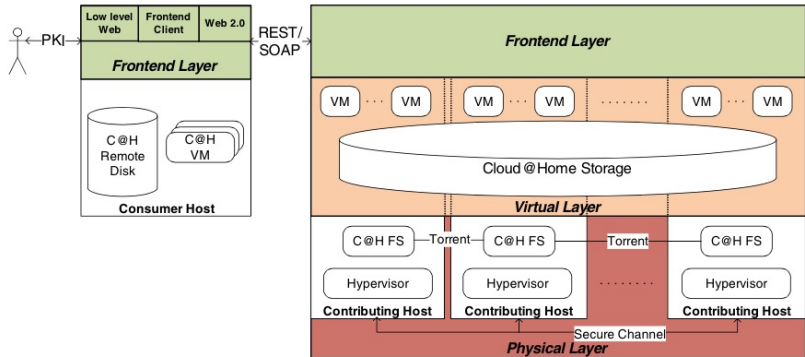
Volunteer Cloud Computing

- Volunteer + Cloud = Volunteer Cloud Computing

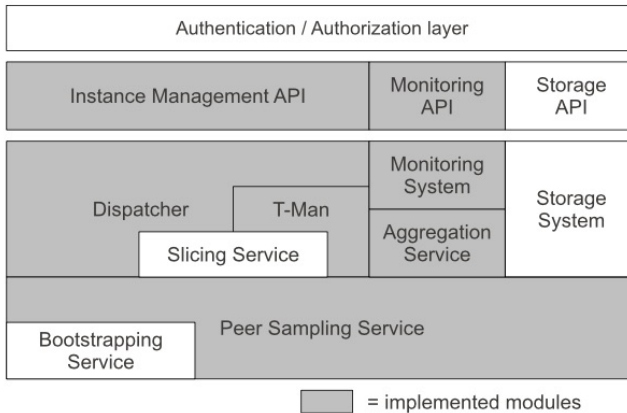
Related Work

- **Cloud@Home**[2009] and **Peer-2-Peer Cloud System**[2011]
- ... and a handful of conceptual reflections

Cloud@Home



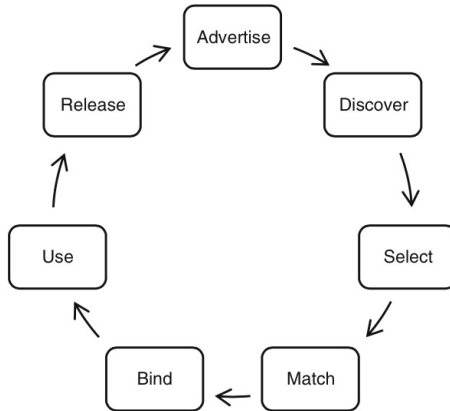
P2PCS



Brief Analysis


- **Scope**
- **Novelty** generally incurs under-specifications of the requirements!

Requirements




Infrastructure


An Architecture for a fully de-centralized peer-to-peer volunteer cloud computing platform-as-a-service infrastructure.

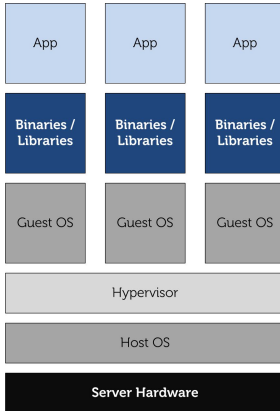


Physical Layer

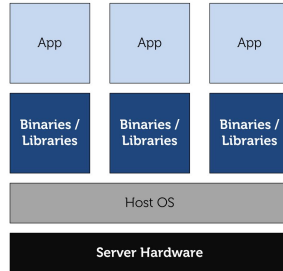


Virtual Layer






Virtualization



Containers

API Layer



- **Structured vs. Unstructured Networks** w.r.t. VCC (trade-off: Single-Attribute-Dominated Queries vs. Multi-Attribute-Dominated Queries).
- **Co-operative Web Hosting** how to host a web application using a peer-to-peer architecture.

- **Using Light Virtualization rather than Traditional (VMs) Virtualization**
- **API Barebone specification**
- **Proof of concept, (work in progress...)**
- **Fully de-centralized approach to VCC**

- Make the proof of concept closer to a production-ready prototype.
- Provide formal analysis of the F.T. and reliability of the system.
- Choose the most adequate solution (Physical Layer), rather than the most convenient.
- Work towards a more concrete ontological representation of VCC.

References I



Ahmed, Reaz and Raouf Boutaba (2014), *Collaborative Web Hosting: Challenges and Research Directions*. Springer.



Babaoglu, Ozalp, Moreno Marzolla, and Michele Tamburini (2012), "Design and implementation of a p2p cloud system." In *Proceedings of the 27th Annual ACM Symposium on Applied Computing*, 412–417, ACM.



Bandara, HMN Dilum and Anura P Jayasumana (2012), "Evaluation of p2p resource discovery architectures using real-life multi-attribute resource and query characteristics." In *Consumer Communications and Networking Conference (CCNC), 2012 IEEE*, 634–639, IEEE.



Bandara, HMN Dilum and Anura P Jayasumana (2013), "Collaborative applications over peer-to-peer systems—challenges and solutions." *Peer-to-Peer Networking and Applications*, 6, 257–276.



Codd, Edgar F (1970), "A relational model of data for large shared data banks." *Communications of the ACM*, 13, 377–387.



Cunsolo, Vincenzo D, Salvatore Distefano, Antonio Puliafito, and Marco Scarpa (2009), "Cloud@ home: Bridging the gap between volunteer and cloud computing." In *Emerging Intelligent Computing Technology and Applications*, 423–432, Springer.



Cunsolo, Vincenzo D, Salvatore Distefano, Antonio Puliafito, and Marco Scarpa (2010a), "Open and interoperable clouds: The cloud@ home way." In *Cloud Computing*, 93–111, Springer.

References II



Cunsolo, Vincenzo D, Salvatore Distefano, Antonio Puliafito, and Marco Scarpa (2010b), "Open and interoperable clouds: The cloud@ home way." In *Cloud Computing*, 93–111, Springer.



De la Rosa, Jose and Kent Baxley (June 2014), "Lxc containers in ubuntu server 14.04 lts."
<http://en.community.dell.com/techcenter/os-applications/w/wiki/6950.lxc-containers-in-ubuntu-server-14-04-lts>. Accessed: 30-12-2014.



Demers, Alan, Dan Greene, Carl Hauser, Wes Irish, John Larson, Scott Shenker, Howard Sturgis, Dan Swinehart, and Doug Terry (1987), "Epidemic algorithms for replicated database maintenance." In *Proceedings of the sixth annual ACM Symposium on Principles of distributed computing*, 1–12, ACM.



Distefano, Salvatore, Maria Fazio, and Antonio Puliafito (2011), "The cloud@ home resource management system." In *Utility and Cloud Computing (UCC), 2011 Fourth IEEE International Conference on*, 122–129, IEEE.



Distefano, Salvatore and Antonio Puliafito (2012), "Cloud@ home: Toward a volunteer cloud." *IT Professional*, 14, 27–31.



Draffan, IW and F Poole (1980), *Distributed data bases*. CUP Archive.



Emelyanov, Pavel and Kir Kolyshkin (2007), "Pid namespaces in the 2.6. 24 kernel." *LWN. net*, November.

References III



Fitzpatrick, Brad (2011), "Memcached: a distributed memory object caching system." *Memcached-a Distributed Memory Object Caching System*.



Ganesan, Prasanna, Beverly Yang, and Hector Garcia-Molina (2004), "One torus to rule them all: multi-dimensional queries in p2p systems." In *Proceedings of the 7th International Workshop on the Web and Databases: colocated with ACM SIGMOD/PODS 2004*, 19–24, ACM.



Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung (2003), "The google file system." In *ACM SIGOPS Operating Systems Review*, volume 37, 29–43, ACM.



Hykes, Solomon (2014), "Dockercon." <http://www.dockercon.com/>. Accessed: 30-12-2014.



Inc., Amazon Web Services (2015), "Amazon web services." <http://aws.amazon.com/>. Accessed: 08-01-2015.



Inc., Google (2014), "Google app engine documentation." <https://cloud.google.com/appengine/docs>. Accessed: 08-01-2015.



Jelasity, Mark and A-M Kermarrec (2006), "Ordered slicing of very large-scale overlay networks." In *Peer-to-Peer Computing, 2006. P2P 2006. Sixth IEEE International Conference on*, 117–124, IEEE.



Jelasity, Márk, Alberto Montresor, and Ozalp Babaoglu (2009), "T-man: Gossip-based fast overlay topology construction." *Computer networks*, 53, 2321–2339.

References IV



Jelasy, Márk, Spyros Voulgaris, Rachid Guerraoui, Anne-Marie Kermarrec, and Maarten Van Steen (2007), "Gossip-based peer sampling." *ACM Transactions on Computer Systems (TOCS)*, 25, 8.



Lewis, Brian (2011), "Separation of responsibilities." Taken from <http://mythoughtsonit.com/2011/04/infrastructure-as-a-service-platform-as-a-service-software-as-a-servicetake-a-look-at-the-management-stack>, Consulted: November 19th 2014.



Licklider, Joseph CR (1963), "Memorandum for members and affiliates of the intergalactic computer network." *Apr*, 23, 350–1.



Linders, James G (1976), "Distributed data bases." *Computers & Geosciences*, 2, 293–297.



Lua, Eng Keong, Jon Crowcroft, Marcelo Pias, Ravi Sharma, Steven Lim, et al. (2005), "A survey and comparison of peer-to-peer overlay network schemes." *IEEE Communications Surveys and Tutorials*, 7, 72–93.



Marzolla, Moreno, Ozalp Babaoglu, and Fabio Panzieri (2011), "Server consolidation in clouds through gossiping." In *World of Wireless, Mobile and Multimedia Networks (WoWMoM), 2011 IEEE International Symposium on a*, 1–6, IEEE.



Maymounkov, Petar and David Mazières (2002), "Kademlia: A peer-to-peer information system based on the xor metric." In *Peer-to-Peer Systems*, 53–65, Springer.

References V



Mell, Peter and Tim Grance (2011), "The nist definition of cloud computing." x.



Menage, Paul (2011), "Linux kernel documentation/cgroups/cgroups. txt."



O'reilly, Tim (2007), "What is web 2.0: Design patterns and business models for the next generation of software." *Communications and Strategies*, 65, 17–37.



Özsu, M Tamer and Patrick Valduriez (2011), *Principles of distributed database systems*. Springer.



Panzieri, Fabio, Ozalp Babaoglu, Stefano Ferretti, Vittorio Ghini, and Moreno Marzolla (2011), *Distributed computing in the 21st century: Some aspects of cloud computing*. Springer.



Pasquet, Mathieu, Francisco Maia, Etienne Rivière, and Valerio Schiavoni (2014), "Autonomous multi-dimensional slicing for large-scale distributed systems." In *Distributed Applications and Interoperable Systems*, 141–155, Springer.



Rimal, Bhaskar Prasad, Eunmi Choi, and Ian Lumb (2009), "A taxonomy and survey of cloud computing systems." In *INC, IMS and IDC, 2009. NCM'09. Fifth International Joint Conference on*, 44–51, Ieee.



Riviere, Etienne and Spyros Voulgaris (2011), "Gossip-based networking for internet-scale distributed systems." In *E-Technologies: Transformation in a Connected World*, 253–284, Springer.

References VI



Vaquero, Luis M, Luis Rodero-Merino, Juan Caceres, and Maik Lindner (2008), "A break in the clouds: towards a cloud definition." *ACM SIGCOMM Computer Communication Review*, 39, 50–55.



Youseff, Lamia, Maria Butrico, and Dilma Da Silva (2008), "Toward a unified ontology of cloud computing." In *Grid Computing Environments Workshop, 2008. GCE'08*, 1–10, IEEE.



Zhang, Qi, Lu Cheng, and Raouf Boutaba (2010), "Cloud computing: state-of-the-art and research challenges." *Journal of internet services and applications*, 1, 7–18.